


Features



References


- <http://szeliski.org/Book/>
- <http://www.cs.cornell.edu/courses/cs5670/2019sp/lectures/lectures.html>
- <http://www.cs.cmu.edu/~16385/>
- <https://medium.com/software-incubator/introduction-to-orb-oriented-fast-and-rotated-brief-4220e8ec40cf>
- https://opencv-python-tutroals.readthedocs.io/en/latest/py_tutorials/py_feature2d/py_table_of_contents_feature2d/py_table_of_contents_feature2d.html
- <https://towardsdatascience.com/sift-scale-invariant-feature-transform-c7233dc60f37>

- Feature detection
 - Talk about harris corner detector?
- Feature descriptor and matching
- Sift
 - Talk about orb
- Panorama


What is a feature?

- There is no universal or exact definition of what constitutes a feature, and the exact definition often depends on the problem or the type of application. Given that, a feature is defined as an "**interesting**" part of an image.
 - [from: wikipedia]


What can we do with features?



Robot navigation



Object recognition




Panorama stitching




Local features: main components


1. Detection: Identify the interest points (also called **keypoints**).



2. Description: Extract vector feature descriptor surrounding each interest point.



3. Matching: Determine correspondence between descriptors in two views



detection

Global detection

- Template matching

Advantages of local keypoints

Locality:

- features are local, so robust to occlusion and clutter

Quantity:

- hundreds or thousands in a single image

Distinctiveness:




- can differentiate a large database of objects

Efficiency

- real-time performance achievable


Local measures of uniqueness

- Suppose we only consider a small window of pixels.
- How does the window change when you shift it?




Local measures of uniqueness


- Suppose we only consider a small window of pixels-
- How does the window change when you shift it?



“flat” region:
no change in all
directions



“edge”:
no change along the edge
direction




“corner”:
significant change in all
directions

Harris corner detection: geometric interpretation

- Consider shifting the window W by (u, v)
 - compare each pixel before and after by summing up the squared differences (SSD).
 - this defines an SSD “error” $E(u, v)$:

$$E(u, v) = \sum_{(x, y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

- **We are happy if this error is high for all $(u, v) \neq (0, 0)$**



Harris corner detection: geometric interpretation

- Taylor Series expansion of I :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

- If the motion (u, v) is small, then first order approximation is good

$$I(x + u, y + v) \approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v$$

- Plug it into the SSD error term:

$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \end{aligned}$$

Harris corner detection: geometric interpretation


$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \\ &\approx Au^2 + 2Buv + Cv^2 \\ &\approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

Also called **second-moments matrix** or **covariance matrix**

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$




Harris corner detection: geometric interpretation

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$


$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Horizontal edge: $I_x = 0$

$$H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$$




Harris corner detection: geometric interpretation

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$


$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Vertical edge: $I_y = 0$

$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$



Harris corner detection: probabilistic interpretation

4. Compute the eigendecomposition

- A real symmetric matrix has an eigendecomposition of:

$$Av = \lambda v$$

$$AQ = Q\Lambda$$

$$A = Q\Lambda Q^{-1}$$

A is real symmetric matrix
 $\xrightarrow{\hspace{1.5cm}}$

$$A = Q\Lambda Q^T$$

$$A = \begin{pmatrix} e_1 & e_2 \end{pmatrix} \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} e_1^T \\ e_2^T \end{pmatrix}$$

– Bonus: eigenvectors are orthonormal if A is real and symmetric.

Harris corner detection: geometric interpretation

- An ellipse can have a matrix form of:

$$x^T (e_1 e_2) \begin{pmatrix} \lambda_1 & \\ & \lambda_2 \end{pmatrix} \begin{pmatrix} e_1^T \\ e_2^T \end{pmatrix} x = 1$$

$$\lambda_1 x^T e_1 e_1^T x + \lambda_2 x^T e_2 e_2^T x = 1$$

$$\frac{(e_1^T x)^2}{\left(\frac{1}{\sqrt{\lambda_1}}\right)^2} + \frac{(e_2^T x)^2}{\left(\frac{1}{\sqrt{\lambda_2}}\right)^2} = 1$$

- Which is exactly as a rotated ellipse with a center of (0,0):



$$\frac{(x \cos(\theta) + y \sin(\theta))^2}{a^2} + \frac{(x \sin(\theta) - y \cos(\theta))^2}{b^2} = 1$$

Harris corner detection: geometric interpretation

- Combining the two equations seen before we can conclude that when taking a cross-section from the error function, we can get an ellipsoid.


$$- [u \ v] H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

- Assume $\lambda_1 > \lambda_2$
- Remember to subtract the mean of each patch so that the ellipsoid is centered.**




Harris corner detection: geometric interpretation

- Eigenvalues and eigenvectors of H
 - e_1 = direction of largest increase in E
 - λ_1 = relative increase in direction e_1
 - e_2 = direction of smallest increase in E
 - λ_2 = relative increase in direction e_2




- λ larger $\leftrightarrow \lambda^{-\frac{1}{2}}$ smaller \leftrightarrow
faster change in $E(u, v)$ in e direction




Interpreting the eigenvalues

- A “good” corner will have a large $R = \lambda_{min}$, which means big change of E in both axis.
- Getting the eigenvectors and eigenvalues is computationally inefficient.
- Instead, use two tricks:
 - $\prod_i \lambda_i = \det(A)$
 - $\sum_i \lambda_i = \text{trace}(A)$
- Then we can more easily compute R :
 - $R = \det(A) - \kappa * \text{trace}(A)^2 \quad (\kappa \in [0.04, 0.06])$
 - $R = \frac{\det(A)}{\text{trace}(A)+\epsilon} = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2 + \epsilon}$


Interpreting the eigenvalues




Interpreting the eigenvalues



I



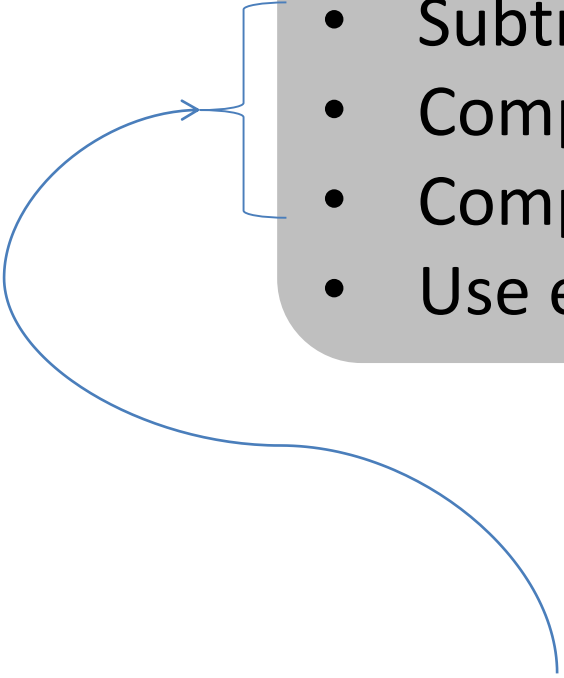
λ_{\max}



λ_{\min}

Harris corner detection: geometric interpretation

- Compute gradients of patch around each pixel.
- Subtract the mean from each patch gradient.
- Compute the covariance matrix.
- Compute eigendecomposition of covariance matrix.
- Use eigenvalues to find corners.




This is PCA (principal component analysis). Out of scope but really interesting!


Harris detector example




f value (red high, blue low)



Threshold ($f > \text{value}$)



Find local maxima of f



Harris features (in red)




Weighting the derivatives

- In practice, using a simple window W doesn't work too well


$$H = \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Instead, we'll *weight* each derivative value based on its distance from the center pixel

$$H = \sum_{(x,y) \in W} w_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



$w_{x,y}$



Gaussian

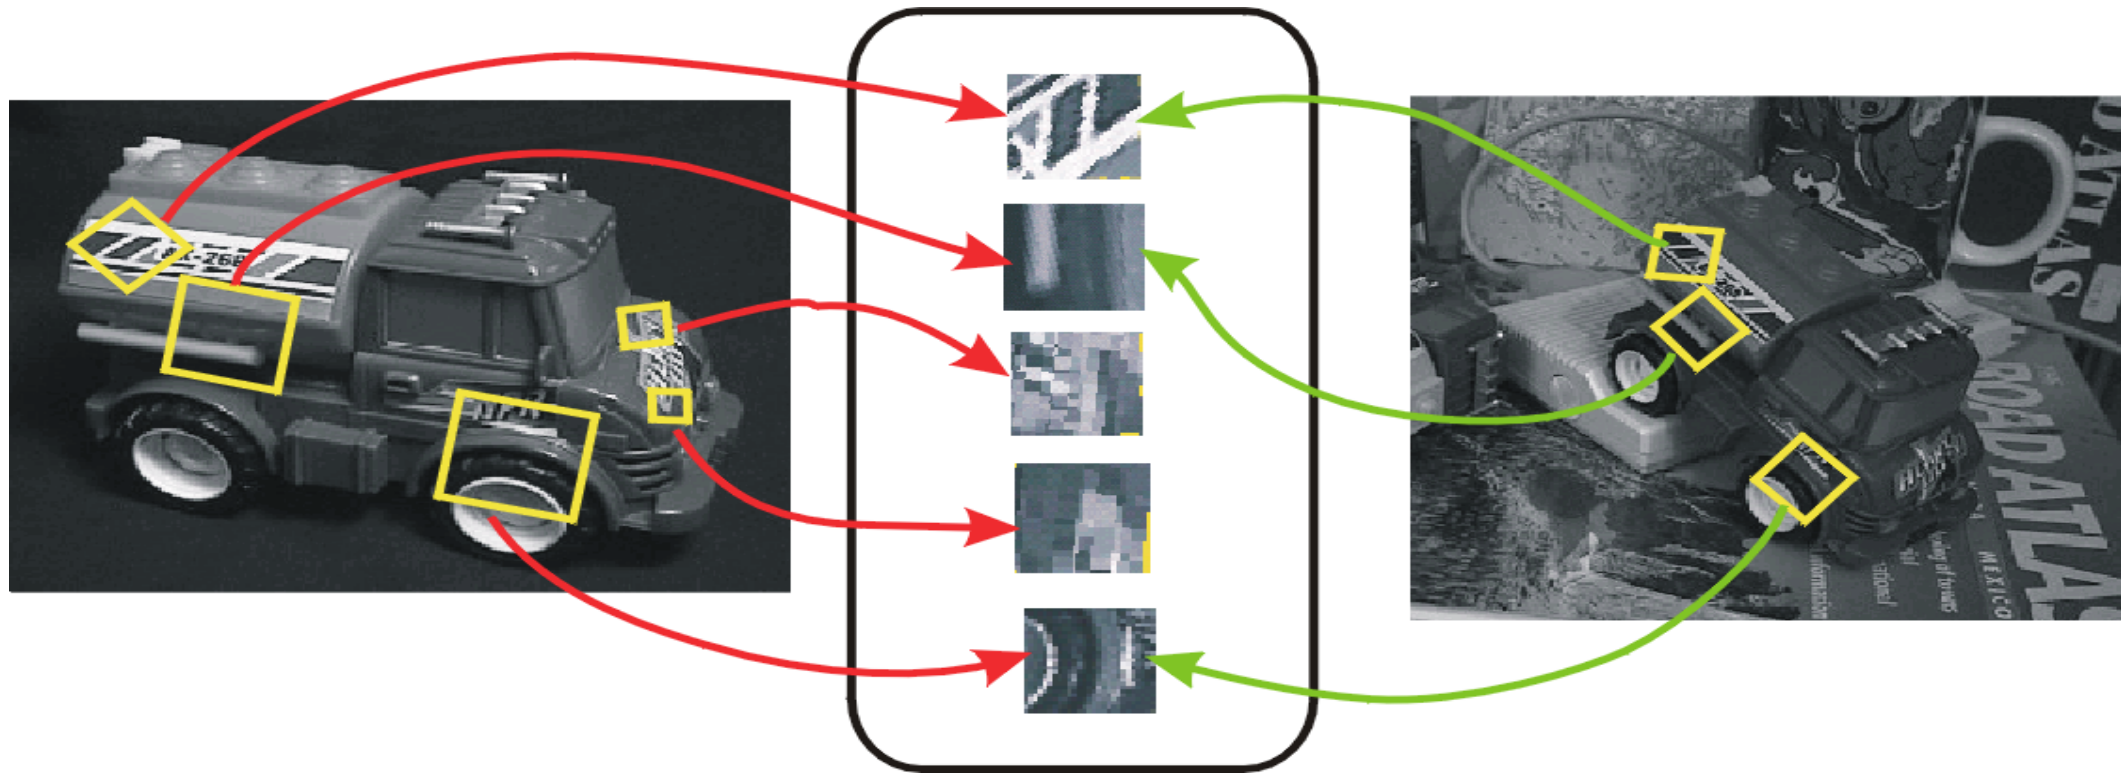


Laplacian

Invariant local features

Find features that are invariant to transformations

- geometric invariance: translation, rotation, scale
- photometric invariance: brightness, exposure, ...



Feature Descriptors


Why extract features?

- Motivation: panorama stitching
 - We have two images – how do we combine them?



Why extract features?

- Motivation: panorama stitching
 - We have two images – how do we combine them?



Step 1: extract features

Step 2: match features

Why extract features?

- Motivation: panorama stitching
 - We have two images – how do we combine them?



Step 1: extract features

Step 2: match features

Step 3: align images